



CZ2005 Operating Systems

Embedded Operating Systems (Nanyang Technological University)

CZ2005 Operating Systems

Introduction

Operating Systems

An OS is a program that acts as a intermediary between a user and computer hardware.

An OS has 2 major goals:

1. User Convenience
2. Efficient hardware utilisation
 - a. Hide hardware complexity
 - b. Use hardware efficiently with smart resource allocation across CPU, I/O, memory

These 2 goals can be contradictory.

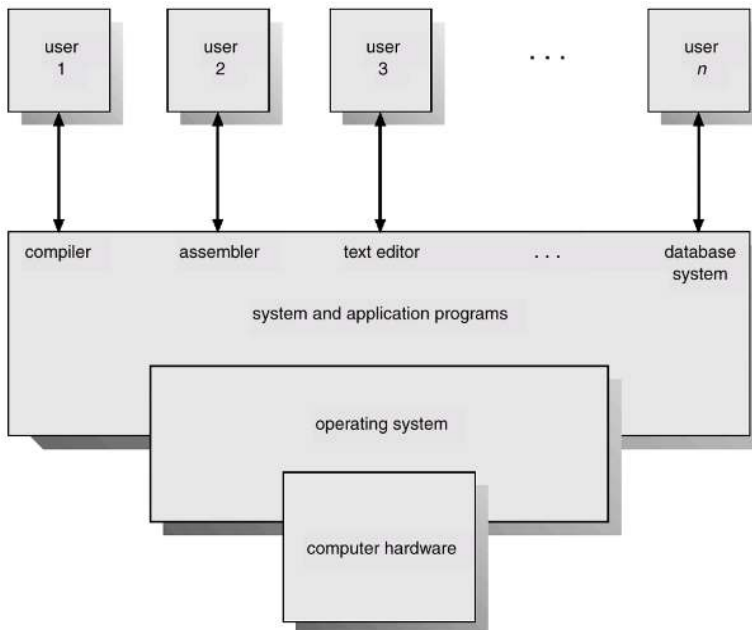
A computer system's components

1. Hardware
2. Operating System
3. Application programs
4. Users

Different Definitions of Operating Systems

1. Resource allocator
2. Control Program
 - a. Controls execution of user programs and operations of I/O devices
3. Kernel
 - a. The 'core' program always ready to accept new commands from users as well as hardware and application programs

An abstract view of system components



Types of computing Systems

1. **Batch systems**
2. **Multiprogrammed and Time-sharing systems**
 - a. E.g: Desktop Systems
3. **Embedded and Cyber-physical systems**
 - a. Real-time systems
 - b. Handheld systems

Simple Batch Systems

- Reduce setup time by batching similar jobs
- Automatic job sequencing, transfers control from one job to another automatically
- Simple memory layout
 - One one user job in memory at a time
- Not very efficient
 - I/O waiting results in idle CPU

Multiprogrammed (Time-sharing) Systems

- Several jobs kept in main memory at the same time, CPU is multiplexed
- Job is swapped in and out of memory
- Highly interactive system that supports multiple online users
 - Desktops, servers

Features needed for Multiprogramming

- Memory Management
 - Allocation across several jobs
- CPU Scheduling
 - To choose among several jobs ready to run
- I/O Device scheduling
 - Allocation of I/O devices to jobs

Desktop Systems

- Personal computers with several I/O devices
 - Keyboard, mouse, display, screen, printers, etc
- Focus on **user convenience** and **responsiveness**
- May run several types of OS

Embedded and Cyber-physical systems

- Physical systems whose operations are monitored and controlled by reliable computing and communication core
- Resource constrained: low power, small memory, low bandwidth etc
- Domain-specific OSes: Real-time, handheld, automotive, avionics, industrial controls, sensor networks, etc

Real-time systems

- Used as a control device in dedicated application in industrial controls, automotive, avionics, medical devices, etc
- Well-defined fixed-time constraints
 - Jobs have a deadline, e.g:airbag control in cars
 - LynxOS, RTLinux, VxWorks Wind River

Handheld Systems

- Mobile phones and tablets
- Issues
 - Limited memory
 - Slow(er) processors
 - Small display screens
 - E.g: iOS, Android, Windows Phone

Multiprocessor Systems

- Systems with more than one CPU, or CPU with multiple cores
- Tightly coupled systems: communications usually through shared memory
- Advantages:
 - Increased system throughput
 - Economical due to sharing of memory and I/O devices compared to single CPU
 - Increased reliability due to redundancy

Computer System Architecture

Computer system operation

- I/O Devices and CPU and execute concurrently
- Device controller with local buffer
- Device controller moves data between buffer and memory
- Informs CPU that it has finished operation causing interrupt

Functions of interrupts

- Interrupt transfers control to **interrupt service routine** through the **interrupt vector**
- Incoming interrupts are disabled while another interrupt is being processed to prevent any loss of interrupts
- A **trap** is a software-generated interrupt caused by error/user request


An operating system is typically interrupt driven

An OS shouldn't poll for tasks/event completions (overhead is high)

Interrupt Handling

- OS preserves state of CPU by storing registers and program counter (context switch)
- OS determines which type of interrupt has occurred
- Based on type of interrupt, identifies the appropriate **interrupt service routine** to execute from the **interrupt vector table**

Interrupt-driven I/O Timeline



Direct Memory Access

- Used for high-speed I/O devices
- OS sets up memory blocks, counters etc
- Device controller transfers data blocks from buffer to main memory without CPU intervention
- One interrupt generated per block rather than per byte

Hardware Protection

Dual Mode Protection

- User mode
- Monitor mode (supervisor/system/kernel mode)
 - Executes OS processes
- Mode bit added to computer hardware to indicate current mode
- Interrupt or fault causes hardware to default to monitor bit
- Privileged instructions only allowed in monitor mode

I/O Protection

- User program may issue illegal I/O operation such as:
 - Read a file that doesn't exist
 - Unauthorised access to device
- ALL I/O instructions are privileged
- All I/O operations must go through OS to ensure correctness and legality
 - Trap is generated for I/O operations trying to bypass OS

Memory Protection

- Interrupt vector and ISR must be protected
- Registers determine range of legal address a program may access
 - Base register: First legal physical memory address
 - Limit register: Size of legal range
- Memory outside defined range protected and inaccessible
- Load instructions of base and limit registers are privileged instructions
- A Load/Store instruction to address:
 - Address checked against base register
 - Address checked against base+limit
 - Trap issued to OS if checks fail

Operating System Services

System Calls

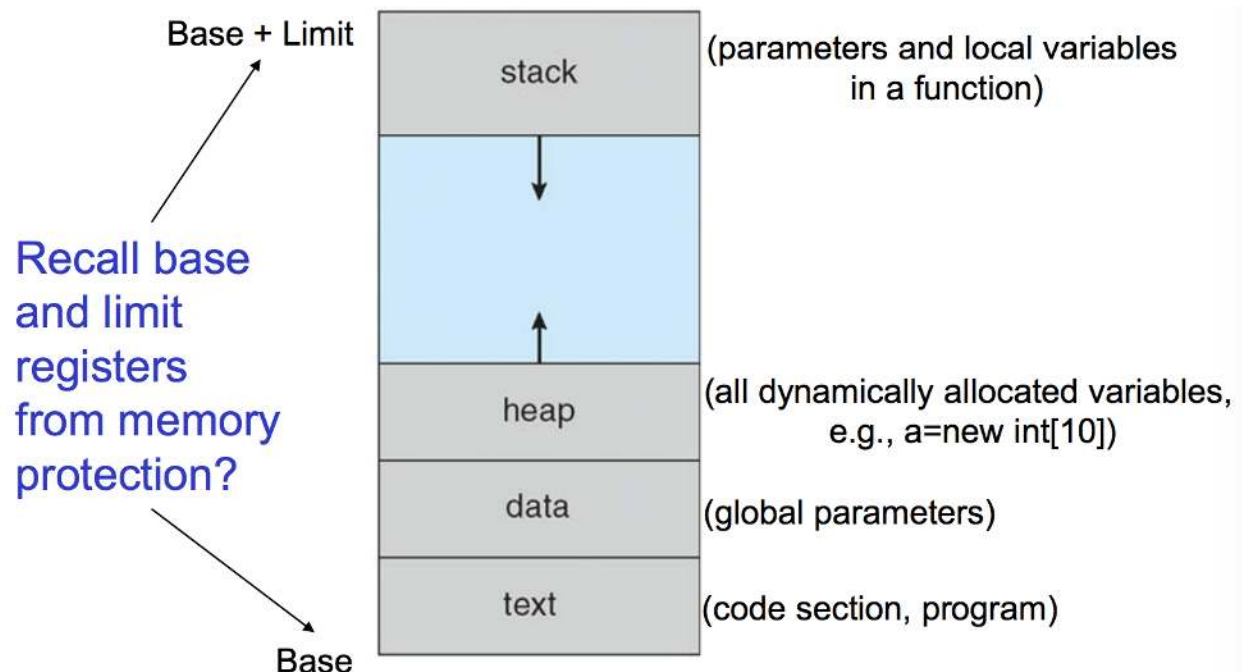
- Interface between user program and operating system (included as standard libraries)
 - Standard libraries act as API for system calls to OS
- Available as assembly-language instructions
- Can be replaced with systems programming to allow system calls made directly (C/C++)
- System call requires switch from user to kernel mode

Processes and Threads

Process Concept

- A process is a program in execution, progressing in sequential fashion
- OS executes variety of process
 - Batch jobs
 - Time-sharing systems (user programs & commands)
- Job = Process

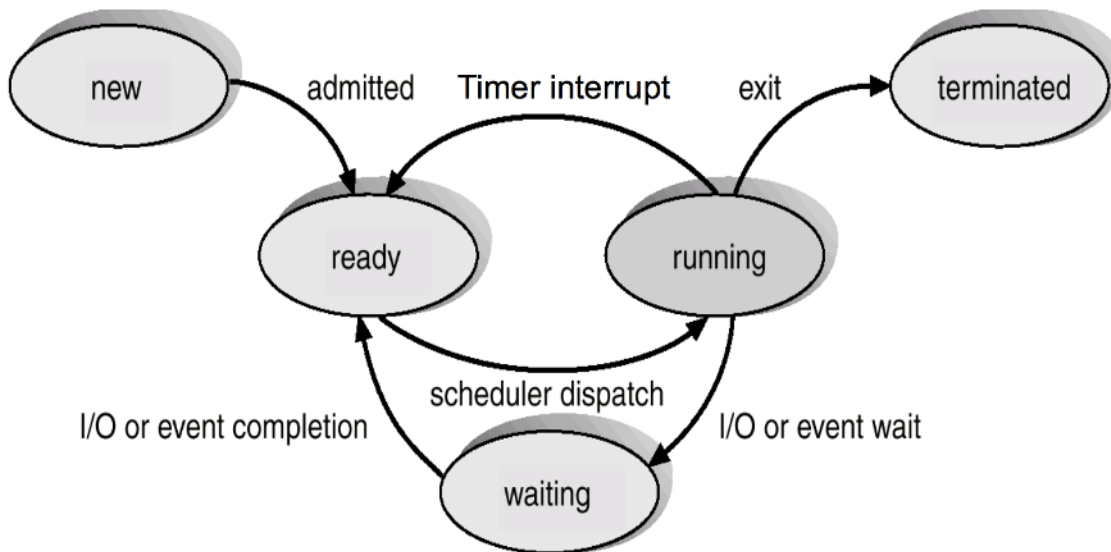
Process in memory



- Global parameters and the program are nearer to base address
- Dynamically allocated variables are on the heap
- Parameters and local variables in a function are in the stack (starting from the limit address of the memory allocated)

Process states

1. New (Process is created)
2. Running (Instruction is executed)
3. Waiting (Process is waiting for an I/O or event)
4. Ready (Process is ready to run and waiting to be assigned to CPU)
5. Terminated (Process has completed)



Timer Interrupt is used to switch between ready processes

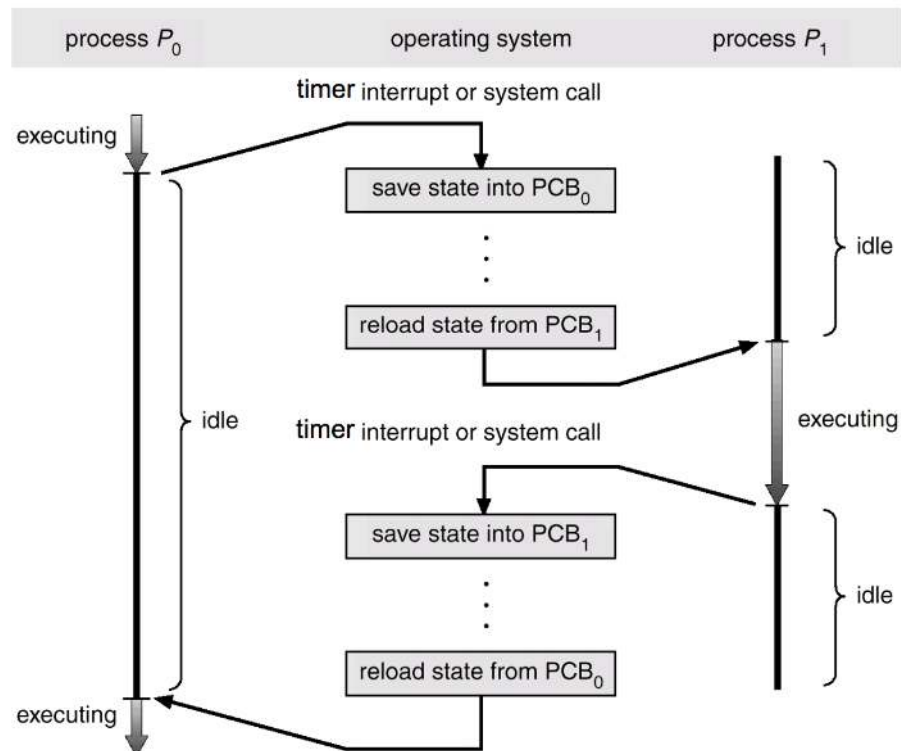
Process Control Block (PCB)

This data structure keeps information associated with each process and stored in kernel memory space. It consists of:

1. Pointer to other PCBs (PCBs maintained in a queue/list structure)
2. Process state
3. Process number (Pid)
4. Program counter (pointer to next instruction)
5. CPU registers
6. Process priority (for scheduling)
7. Memory management information (base and limit register values)
8. Information regarding files (list of open files)
9. Pointer to next PCB

Process scheduling

Context Switch between processes



When a CPU switches to another process, the system must save the context of the old process and load the saved context for the new process.

Context switch time is overhead. System does no useful work while switching.

Process scheduling queues

Job queues: Set of all processes with the same state

- Ready queue

- Device queue (processes in waiting state waiting for specific I/O device)

All queues are stored in main memory (kernel memory space)

Processes migrate between queues when state changes

Process schedulers

1. Long-term scheduler (job scheduler)
 - a. Selects process from disk and loads them into main memory for execution
 2. Short-term scheduler (CPU scheduler)
 - a. Selects from processes that are ready to execute and allocates CPU to one of them
 3. Medium-term scheduler
 - a. When system load is heavy, swaps out partially executed process from memory to hard disk
 - b. When system load is lighter, these processes can be swapped back into main memory
- 2)** is invoked frequently (100ms) in multiprogrammed system for responsiveness and efficiency purposes
- 1)** is invoked infrequently (seconds/minutes)

Degree of multiprogramming is initially controlled by long-term scheduler and *thereafter* controlled by medium-term scheduler.

Process creation

Parent processes create children processes and so on, forming a tree of processes

This can lead to two possible execution orders:

1. Parent and child execute concurrently and independently
2. Parent waits until all children terminate (`wait()`, `join()`)
 - a. E.g: Web browsers nowadays fork a new process when making new tab
 - b. OS creates background processes for monitoring and maintenance

Process Termination

1. Exit: process executes last statement and asks OS to delete it
 - a. Child output return data to parent
 - b. Process resources are de-allocated by OS
2. Abort: Parent terminate execution of children processes at any time
 - a. Child exceeded allocated resources
 - b. Task assigned to child is no longer required
 - c. Parent is exiting

Cooperating Processes

- Independent process cannot affect or be affected by execution of other process
- Cooperating processes can affect or be affected
 - Such processes have to communicate with each other to share data
 - Inter-Process Communication
 - Message Passing
 - Shared Memory
- **Message passing:**
 - No shared variables
 - 2 system calls needed:
 - send(message, recipient/mailbox)
 - receive(message, sender/mailbox)
 - Mailboxes are used for indirect message passing (AKA ports)
 - Communication requires a communication link and use system calls

Producer-Consumer Process Paradigm

Classical paradigm for cooperating processes

- Producer produces information for consumer that uses it
- Shared buffer used for storing information
 - Unbounded or bounded buffer size
- Producer waits while mailbox is full. Consumer waits while mailbox is empty.

Process Scheduling

Basic Concepts

CPU, I/O Bursts: Process execution alternates between CPU and I/O operations

CPU Burst: Duration of one CPU execution cycle (multiple ticks of a CPU)

I/O Burst: Duration of one I/O operation (wait time)

CPU Scheduling objective: Keep CPU busy as much as possible (multiprogramming)

Focus: short-term (ready queue) scheduler on single CPU

CPU Scheduler types:

1. Nonpreemptive: Once CPU has been allocated a process, process will keep CPU until it voluntarily releases CPU by termination or requesting I/O or event wait.
2. Preemptive: CPU can be taken away from running process at any time by scheduler
 - a. Through interrupt, during process creation, or after I/O or event completion

Scheduling Objectives

1. Max CPU Utilisation (CPU busy as possible)
2. Max throughput (Number of processes that complete their execution per unit time)
3. Min turnaround time (time of creation to termination)
4. Min waiting time (amount of time process waits in ready queue)
 - a. $\text{Waiting time} = \text{turnaround time} - \text{CPU burst}$ (if process have single CPU burst with no I/O)
5. Min response time (Amount of time from request submission to create process until first response is produced)

Scheduling algorithms

1. First-come, first-served (FCFS)
2. Shortest Job First
3. Priority-based Scheduling
4. Round-robin
5. Multilevel queue scheduling

Gantt chart execution for CPU using FCFS



Waiting time: 0, 24, 27

Average waiting time: 17 seconds

If arrival order is changed to P₂, P₃, P₁, then average waiting time is reduced to 3 seconds

FCFS Properties:

1. Nonpreemptive
2. Suffers from convoy effect due to earlier arrived long process

SJF Scheduling: Prioritize processes based on CPU burst lengths

SJF/SRTF is **OPTIMAL** for minimising average waiting time.

When SRTF preempts a process out of the CPU, the waiting time still increases (since it is placed back into the ready queue).

Priority-based Scheduling:

Priority number is assigned to process with highest priority

- Smallest integer = highest priority
- Two schemes: Preemptive and nonpreemptive
- FCFS = priority based on arrival order
- SJF = priority based on CPU burst length
- **Starvation issue: Lower priority processes may never execute in heavily loaded system**
 - **Solution: ageing. Increase priority of processes that are not able to execute**

Round-robin Scheduling

Fixed time quantum for scheduling (q): Process is allocated CPU for q time units, preempted thereafter and inserted at the end of the ready queue. q is usually around 10-100ms

Performance:

- N processes in ready queue: waiting time no more than $(n-1)q$ time units
- Large q = degenerates to FCFS
- Small q = Too many context switches, high overhead
- **Generally higher waiting/turnaround time than SFJ**
- **Better response time**

Multilevel Queue scheduling:

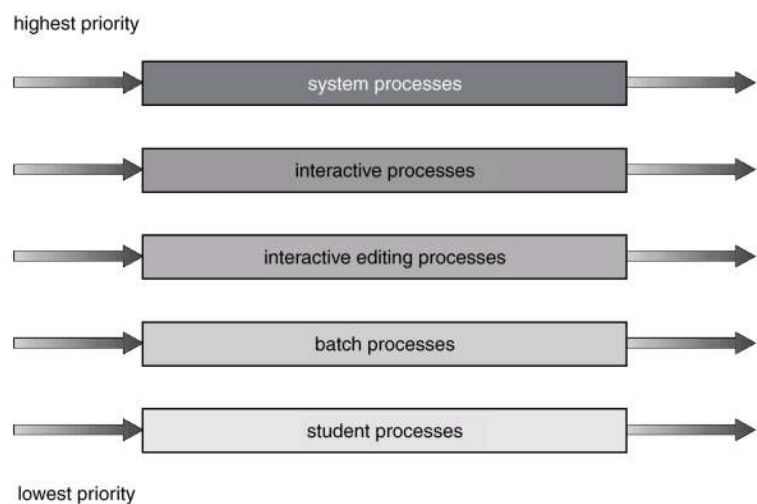
Different processes have different requirements

- Foreground processes like I/O handlers need to be interactive (RR preferred)
- Background processes need not be interactive, scheduling overhead low (FCFS)

In MLQ, Ready queue is partitioned into several queues with its own scheduling algorithm

Two schemes allowing for inter-queue scheduling:

1. Fixed priority scheduling: queues served in priority order (foreground before background)
 - a. **Starvation issue**
2. Time-slice based scheduling: Each queue gets fixed time quantum on CPU
 - a. 80ms foreground, 20ms background



Process synchronisation

Race Condition & Critical Section

N processes competing to use shared data that requires a proper order of execution for concurrent processes. Causal ordering of reads and writes.

Critical section problem:

Each process has a code segment called a critical section, in which shared data is accessed. At least one process modifies the shared data.

A protocol is required to ensure that when one process is executing in critical section, no other process is allowed to execute in critical section.

Assumption of critical and remainder sections:

1. Each process is guaranteed to make progress over time in these sections
2. Execution speed may be different

Desired Properties of entry and exit sections that bound the critical section:

1. **Mutual exclusion:**
If a process is executing in its critical section, then no other process can be executing in its own critical section at the same time.
2. **Progress:**
If no process is executing in its critical section and there exists processes that wish to enter their critical section, then the selection of the next process to enter the critical section cannot be postponed indefinitely.
3. **Bounded waiting:**
After a process has requested to enter critical section, and before that request is granted, other processes are allowed to enter their critical section only a bounded number of times.

User-level software approaches w/o OS support

Property Analysis	Algorithm 1 (turn)	Algorithm 2(flag)	Algorithm 3 (flag & turn)
Shared variables	<pre>int turn = 0;</pre> <p>turn = i means P_i can enter critical section</p>	<pre>boolean flag[2]; flag[0]=flag[1]=false;</pre> <p>Flag[i] = true means P_i ready to enter critical section</p>	<pre>int turn = 0; boolean flag[2]; flag[0]=flag[1]=false; Flag[i] = true means P_i ready to enter critical section turn = i means P_i can enter critical section</pre>
Entry section	<pre>while(turn != i);</pre>	<pre>flag[i] = true; while (flag[k]);</pre>	<pre>flag [i] = true; turn = k; while (flag [k] and turn = k);</pre>
Exit section	<pre>turn = k;</pre>	<pre>flag[i] = false;</pre>	<pre>flag [i] = false;</pre>
Mutual Exclusion	Yes	Yes	Yes
Progress	<p>No</p> <p>Turn = 0 and P_0 is in a long remainder section, and P_1 wants to enter critical section Turn is not returned to P_1</p>	<p>No</p> <p>P_0 executes flag [0] = true; Context switch to P_1 P_1 executes flag [1] = true; Both are stuck in infinite while loop</p>	Yes
Bounded Waiting	Yes	Yes	Yes

OS-level solutions

Software approaches are difficult to implement for more than 2 processes

OS support has 3 kinds:

1. Synchronisation hardware
2. Semaphore
3. Monitor

Modern processors provide special atomic hardware instructions (non-interruptible)

TestAndSet: Test and modify content of main memory word atomically

```
boolean TestAndSet(boolean *target){
    boolean rv= *target;
    *target = true;
    return rv;
}
```


Property Analysis of TestAndSet:

1. **Mutual Exclusion:** Yes
2. **Progress:** Yes
3. **Bounded Waiting:** No.

Proof:

• Process P_0

1. Initialization code

2. while (1) {

3. while(TestAndSet(&lock));

4. critical section

5. lock = false;

6. }

• Process P_1

1. while (1) {

2. while(TestAndSet(&lock));

3. critical section

4. lock = false;

5. }

Context switch points

- Suppose for P_0 lines 1-3 take 10ms and 2-6 also take 10ms
- Suppose we use round-robin scheduling with quantum 10ms
- Whenever P_0 context switches out, it holds the lock (is in critical section)



Semaphore

Shared integer variable.

Accessible by two atomic operations:

1. Wait(S): if $S > 0$, $S--$; else, yield();
 - a. yield(): enter waiting state or busy loop waiting
2. Signal(S): $S++$;

Two kinds of semaphores:

- 1) Counting semaphore: Integer value unrestricted
- 2) Binary semaphore: Value can never be greater than 1

Classical Semaphore Implementation

Wait(S): while($S \leq 0$); $S--$;

Signal(S): $S++$;

Pros: No context switch overhead (busy waiting)

Cons: Inefficient if critical sections are long, busy waiting wastes CPU cycles.

Mutual Exclusion: Wait(S) must be atomic. No context switch allowed between while and $S--$;
If not, mutual exclusion fails.

Current Semaphore Implementation:

```
typedef struct {  
    int value;  
    struct process *L;  
} semaphore;
```

L is a queue that stores processes waiting on this semaphore in a form of a PCB list

`block()` : Dequeue current process from ready queue

Enqueue process to L

Change state of process to waiting

`wakeup()` : Dequeue current process from L

Enqueue process to ready queue

Change state of process to ready

```
Wait(S) : S.value--;  
    If (S.value < 0) {  
        block();  
    }
```

```
Signal(S) : S.value(++)  
    If (S.value <= 0) {  
        wakeup(P);  
    }
```

Best practices for semaphores

1. Find shared variables and data/shared resources
2. Identify critical section
3. Protect shared variables
 - a. Use binary semaphore for mutex
 - b. Apply wait and signal
4. Protect shared resource
 - a. Use one counting semaphore per resource
 - b. Initial value = number of resource instances
 - c. Apply wait and signal

Starvation

Priority based or LIFO policy queuing can result in starvation, as a process under these policies may never have a chance to be removed from semaphore queue.

Classical Synchronisation Problems

Bounded-buffer (Producer-Consumer) Problem

Multiple consumer/producer environment

Buffer is a variable: apply binary semaphore (mutex)

Full and empty are used as resource counting semaphores

Full initially 0, empty initially n

First wait on empty/full (producer/consumer)

Then request for mutex

After critical section

Release mutex

Signal full/empty (producer/consumer)

Dining Philosophers Problem

Five resources, Five processes, Each needs 2

Solutions: Allow at most 4 philosophers to be hungry simultaneously

Allow a philosopher to pick up his chopsticks only if both chopsticks are available

Use an asymmetric solution: Odd philosopher picks up left chopstick first then his right chopstick

Even philosopher picks up right chopstick then left chopstick

Readers-Writers Problem

Writer requires exclusive access

Multiple readers can concurrently access

First reader: Block writers

Last reader: After reading, can allow writers

Shared data: readcount=0;

Semaphores: readCounter = 1; databaseAccess = 1;

Writer: wait(databaseAccess);
 signal(databaseAccess);

Reader: wait(readCounter); readcount++;
 If (readcount == 1) wait(databaseAccess);
 signal(readCounter);

 wait(readCounter);
 readcount--;
 if(readcount = 0) signal(databaseAccess);
 signal(readCounter);

Deadlocks and Starvation

The Deadlock Problem

Deadlock: A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set.

The shared resource deadlock can be resolved if one car backs up (release resource and rollback)

Several processes may need to rollback if a deadlock occurs

Resource Allocation Graph

Graph of vertices V is partitioned into two types:

$P = \{P_1, P_2, \dots, P_n\}$, consisting of all processes in system

$R = \{R_1, R_2, \dots, R_m\}$, consisting of all resource types in the system


Edges E is partitioned into two types:

Request edge: Directed edge $P_i \rightarrow R_j$

If request is granted, edge is removed


Assignment edge: Directed edge $R_j \rightarrow P_i$

If resource is released, edge is removed


Process  Resource Type with 4 instances



P_i requests instance of R_j



P_i is holding an instance of R_j



If the graph has no cycles, there is no deadlock

If the graph contains cycles, and there is only one instance per resource type, then there is deadlock. If several instances exist per resource type, there is a possibility of deadlock.

Deadlock conditions

Mutual exclusion: Only one process can use a resource at a time.

Hold and wait: A process holding at least one resource is waiting to acquire additional resources held by other processes

No preemption: A resource can be released only voluntarily by the process holding it, after that process has completed its task

Circular wait: There exists a set of processes of processes waiting for resources such that there is a cycle in waiting for resources within the set.

Handling deadlocks

Ensure system will never enter deadlock state

Allow system to enter deadlock and then recover

Ignore problem and pretend deadlocks never occur

Prevent at least one of the four deadlock conditions

Deadlock Avoidance Algorithm

Algorithm checks for safe state of resource allocation.

Safe state is when there exists safe completion sequence of all processes without deadlock

- 1) Find a process P_1 where $\text{Request} \leq \text{Available}$
- 2) Find a process P_2 where $\text{Request} \leq \text{Available} + P_1.\text{Hold}$
- 3) Find a process P_3 where $\text{Request} \leq \text{Available} + P_1.\text{Hold} + P_2.\text{Hold}$
- 4)
- 5) If all processes can be included, the sequence P_1, P_2, P_3, \dots is safe.

System in unsafe state means there is a **possibility of deadlock**. If a process releases resources before its completion then deadlock may not occur

Banker's algorithm

Algorithm to check whether satisfying a resource request would lead to unsafe state.

Assumptions:

1. Each process must declare the maximum instances of each resource type that it needs
2. When a process gets all its resources it must return them in a finite amount of time

Data structures, where m = number of resource types, n is number of processes

1. available (vector of length m) where if $\text{available}[j] = k$, there are k instances of resource
2. Max ($n \times m$ matrix): if $\text{Max}[i,j] = k$, process P_i can request at most k of R_j
3. Allocation ($n \times m$ matrix): similar to above but about current allocation
4. Need ($n \times m$ matrix): P_i may need k more instances of R_j to complete its task

Banker's algorithm part 1: Safety algorithm

1. Work and finish are vectors of length m and n
2. Work = available, finish[i] = false for all i
3. Find i such that Finish[i] == false and Need[i,*] < work
4. If no such i exists, go to step 6
5. Work = work + allocation[i, *]; Finish[i] = true
6. If finish[i] == true for all i, system is safe. Else, system is unsafe.

Part 2: Resource request algorithm

Request vector from P_i

- 1) If Request \leq Need_i go to 2). Else, raise error condition
- 2) If Request \leq Available go to 3). Else, wait
- 3) Pretend to allocate requested resources to P_i
 - a) Available = Available - request
 - b) Allocation = Allocation + request
 - c) Need = need - request
- 4) Run safety algorithm to check if safe.
- 5) If safe, allocate resource. Else, unsafe, P_i must wait and state is restored to pre 3).

Memory organisation

Binding code and data to memory

To run a program, a process image must be created and loaded into memory

Instructions in memory must be fetched to the CPU. Instruction may require another access to memory.

Address binding of instructions and data to memory address can happen at 3 stages:


1. Compile time: If memory location is prior known, absolute code can be generated. Recompile required if starting location changes
2. Load time: Compiler generates relocatable code. Binding performed by image loader.
3. Execution time. If process can be moved during execution from one memory segment to another, binding is delayed until runtime.

In compile- or load-time binding, absolute address format is used. If address generated by CPU $< \text{base}$ or $\geq \text{base} + \text{limit}$ → error
Otherwise, memory address = address generated by CPU

In execution-time binding, relative address format is used. If address generated $\geq \text{limit}$ → error.
Otherwise, memory address = address generated by CPU + base

Logical vs. Physical Address Space

- Logical address: generated by CPU; also referred to as virtual address
- Physical address: address seen by memory unit
- Address space: Address accessible by process



Memory allocation among processes

- Contiguous allocation: Logical address space of process remains contiguous in physical memory
 - Fixed partitioning
 - Dynamic partitioning
- Non-contiguous allocation: A process logical address space is scattered over different regions in physical memory

Fixed Partitioning

Memory is partitioned into regions with fixed boundaries

Contiguous Allocation

Hole: block of available memory; holes of various size are scattered throughout memory

When a process arrives, it is allocated memory from a hole large enough to accommodate it.

Operating system maintains information about allocated partitions and free partitions (holes)

Dynamic Storage Allocation Problem

How to satisfy a request of size n from list of free holes?

- First-fit: First hole big enough
- Best-fit: Smallest hole big enough. Must search entire list unless order by size. Produces smallest leftover hole.
- Worst-fit: Allocate largest hole; must also search entire list. Produces largest leftover hole.

Fragmentation:

- External fragmentation: Total memory space exists to satisfy a request but it is not contiguous. Happens outside partitions
 - Reduce via compaction: Shuffle memory contents to place all free memory together in one large block.
 - Possible only if relocatable address format is used in process image and binding is done during execution time.
- Internal fragmentation: Allocated memory may be slightly larger than requested memory, size difference is memory internal to partition, but not being used.

Paging

Memory space allocated to process can be noncontiguous, process is allocated physical memory whenever latter is available.

Divide physical memory into fixed-size blocks called frames(size is power of 2, 512-8192 bytes)

Divide logical memory of same size called pages.

- Keep track of all free frames
- Setup page table to translate logical to physical addresses
- External fragmentation is eliminated
- Internal fragmentation is still possible, on average half of a page is unused.

Address Translation Scheme

Logical address contains:


- Page number (p) used as index into page table entry containing frame number in physical memory
- Page offset (d) combined with frame number to define physical memory address that is sent to memory unit

Page size: 2^n

Logical address space: 2^m

Number of pages: 2^{m-n}

Page table entry size:



Logical address: p, d: E.g: 0111

Since we know that the logical address' last n bits are the page size (i.e page offset)

So the other bits are the index for the page table to find the frame number in physical memory.

Thus, p = 01, d = 11.

P corresponds to index 1 (page 1) on page table.


Page table indicates frame 6 in main memory.

6 = 110, including offset d, the physical address will be 11011 → 27 in main memory

Implementation of page table

- Kept in main memory
- Has 2 registers assigned: Page-table base register and Page-table length register
- Every data/instruction access requires two memory accesses: one for page table and one for data/instruction - effective memory access time: $2t$ with t being memory cycle time unit
- Memory access time can be reduced by use of fast-lookup hardware cache
 - Associative registers (think register array)
 - translation look-aside buffers

Paging using TLBs



TLB Lookup = l time unit

Memory cycle time = t time unit

Hit ratio: percentage of times page number is found in TLB: r

$$\begin{aligned}\text{Effective access time} &= (t + l)r + (2t + l)(1-r) \\ &= (2-r)t + l\end{aligned}$$


Two-level Page-Table Scheme

A logical address (32-bit machine, 4K page size)

2^{20} entries in a page table

If each entry is 4 bytes, each table occupies 4 megabytes of memory

A large page table is divided up to be easier to allocate in main memory, with small increase in effective access time



Logical address on 32-bit 4K page size machine has:

- Page number of 20 bits
- Page offset of 12 bits

Page table is paged: Page number divided into:

- A 10-bit p_1 : outer page table contains $(4 \times 2^{20}) / 4K = 2^{10}$ entries
- A 10-bit p_2 : each page in page table contains $(4 \times 2^{10}) / 4 = 2^{10}$ entries

Inverted Page Table

Usually each process has its own page table. System could have many page tables consuming substantial memory space

Page table size is proportional to logical address space

Alternative: have single table with one entry for each physical frame, as $\langle \text{process-id}, \text{page-no} \rangle$. This is an Inverted Page Table

Logical address: $\langle \text{process-id}, \text{page-no}, \text{offset} \rangle$

Increases search time: table sorted by physical address but lookups occur on logical address

Inverted as the table is sorted by frame number not page number.

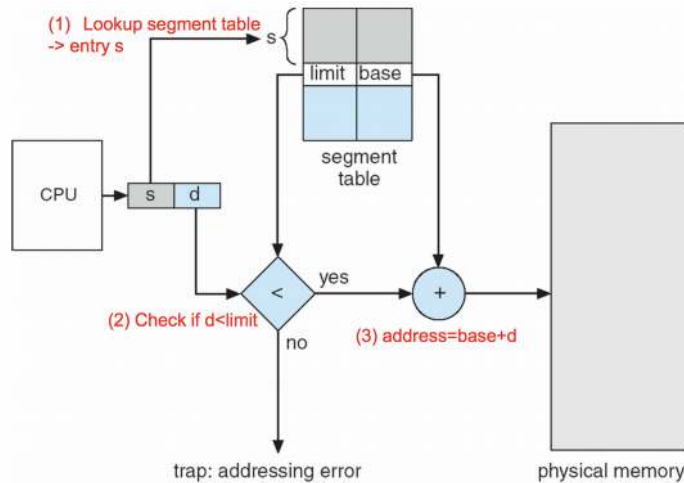
When searching, use $\text{pid} + \text{page number}$ to find which frame it is in. Then physical address is the frame number + page offset

Segmentation

- Memory-management scheme that breaks program up into logical segments and allocates space for these segments into memory separately
- Unlike pages, segments are variable size
- A process has collection of segments
- Like pages, segments may not be allocated contiguously

Address Translation

- Each segment has segment no. & offset.
- Segment table. Each table entry has:
 - Base - starting physical address
 - Limit - length of segment
- Segment-table base register points to segment table's location in memory
- Segment-table length register indicates number of segments used by program



Fragmentation in segmentation

- With variable segment length, external fragmentation is a problem. Usually use best-fit or first-fit. Dynamic storage-allocation problem.

Virtual memory

Swapping

- A process can be swapped out temporarily out of memory to backing store
- Backing store: Fast disk large enough to accomodate copies of all memory images for all users; must provide direct access to these memory images
- Major part of swap time is transfer time: total transfer time is directly proportional to amount of memory swapped
- Swapping used to
 - make space available for processes that require more memory
 - Increase degree of multiprogramming (medium-term scheduler)

Assumption was that whole program was loaded into memory at some point of time.

This restricts program size to size of main memory.

To remove this restriction, logical size requirement must be decoupled from size of physical memory space

Demand paging

- Implementation of virtual memory is implemented using demand paging
 - With each page table entry there is a valid-invalid bit
 - Initially all invalid
 - During address translation, if bit is invalid, there is page fault
 - Page is brought into memory
1. CPU reference to page 6
 2. Page table checked for page 6: invalid
 - a. State of process is saved when trap to OS activated
 3. Page is on backing store
 4. Page is brought into memory from disk
 5. Page is placed in physical memory in free frame
 6. Page table is reset
 - a. Valid-invalid bit set to valid
 - b. Page table entry updated
 7. Restart instruction 1)

Performance of Demand Paging

p = probability of page fault

$EAT = (1-p) * \text{memory access time} + p * \text{page fault time}$

Memory access time \ll page fault time

Page fault time:

1. Servicing page fault interrupt
2. Read new page (longest time)
3. Restart process

Page Replacement

If there's no free frame, a page not in use must be paged out

Page transfer may be necessary (one out one in) and increases page fault time

1. Find victim page using place replacement algorithm
 - a. Check dirty/modified bit, if true, page out
2. Change page table entry of page table to invalid
3. Bring desired page into freed frame
4. Update the page table entry for the new page

Page-replacement algorithm

Objective: Algorithm with lowest page-fault rate

Evaluation is through a string of memory references called a reference string

Count page faults

As number of available frames increase, number of page faults decrease

Algorithms

FIFO Algorithm

Replace page that has been loaded in memory for longest time

Iterate over array and wrap around

Belady's Anomaly

Increasing number of frames lead to more page faults

	1	2	3	4	1	2	5	1	2	3	4	5
F1	1	1	1	4	4	4	5	5	5	5	5	5
F2		2	2	2	1	1	1	1	1	3	3	3
F3			3	3	3	2	2	2	2	2	4	4
	P	P	P	P	P	P	P			P	P	

	1	2	3	4	1	2	5	1	2	3	4	5
F1	1	1	1	1	1	1	5	5	5	5	4	4
F2		2	2	2	2	2	2	1	1	1	1	5
F3			3	3	3	3	3	3	2	2	2	2
F4				4	4	4	4	4	4	3	3	3
	P	P	P	P			P	P	P	P	P	P

9 page faults vs 10 page faults

Optimal Algorithm

Replace page that will be used for the longest period of time:

Used as theoretical benchmark.

Inclusion property: Pages loaded in n frames is always a subset of pages in $n+1$ frames

Least Recently Used (LRU) Algorithm

Replace page that has not been used for the longest period of time

1. Counter implementation
 - a. Each page table entry has last-used entry
 - b. CPU increments logical clock for each memory reference
 - c. Whenever a page is referenced, clock value to copied to last-used field
 - d. During page fault, page with smallest last-used value is kicked out
2. Stack implementation
 - a. Keep a stack of page numbers in doubly linked list
 - b. Referenced page is moved to the top
 - c. Bottom of stack is page to page out
 - d. List update is expensive, but no search needed for replacement

LRU Approximation algorithms

Hardware support necessary for exact algorithms are expensive and generally unavailable


Hardware settable reference bit is used

1. Each page has a bit R initially 0
2. On page reference, R is set to 1
3. Replace page with R = 0 if exists.

Second Chance Algorithm

Clock Algorithm:

1. Clock hand points to oldest page
2. During page fault, page being pointed by hand is inspected. If R bit is set, it is cleared and the hand is advanced to the next page
3. Otherwise page is evicted and new page is inserted into its place, hand is advanced one position
4. Worst case: when all bits are set, hand cycles through whole queue giving each page a second chance. Degenerates to FIFO replacement



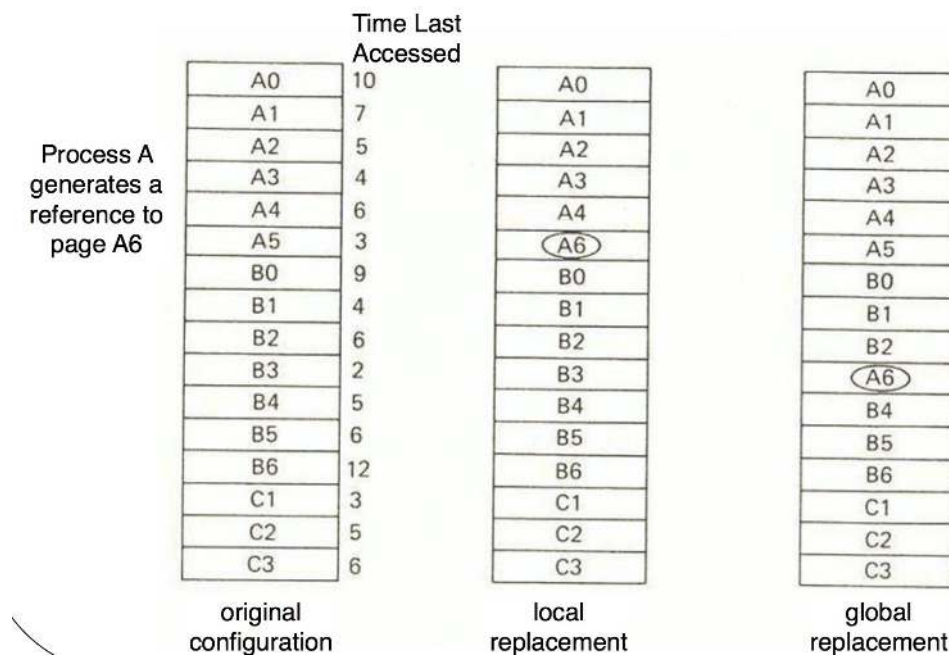
Allocation of Frames

- Smaller memory allocated, more processes residing in memory
 - Small number of pages loaded increases page faults
 - Past a certain size, further allocation of pages will not affect page fault rate
1. Fixed allocation
 - a. Fixed number of frames (equally or proportionally allocated)
 - b. During page fault, one of the pages of that process must be replaced
 2. Variable allocation
 - a. Allows page frame allocated to a process be varied over process lifetime

Global vs Local Allocation

Global Replacement: Process selects replacement frame from set of all frames; one process can take a frame from another process. Performance dependant on external circumstances

Local Replacement: Each process selects only from its own set of allocated frames. May hinder other processes by not making available its less used pages/frames



	Local Replacement	Global Replacement
Fixed Allocation	<ul style="list-style-type: none"> Number of frames allocated is fixed Page to be replaced is chosen from frames allocated to process 	Not possible
Variable Allocation	<ul style="list-style-type: none"> Number of frames allocated variable across process lifetime Page to be replaced is chosen from frames allocated to process 	Page to be replaced chosen from all available frames in main memory

Thrashing

If a process does not have enough pages, page-fault rate is high. Leads to:

1. Low CPU utilisation
2. OS thinks degree of multiprogramming needs to increase
3. Another process is added
4. CPU utilisation drops further

Working-set Model

Locality of reference:

- Temporal Locality: Locations referred to just before are likely to be referred to again. \
- Spatial Locality: Code & data are usually clustered physically

Working-set window: fixed number of page references

Working set is considered as the set of unique frame references in the last n references

WSS_i (working set size of Process P_i): total number of pages referenced in most recent working set window


$D = \sum WSS_i$: total demand for frames

M : total number of frames

If $D > m$: Thrashing

Policy: if $D > m$, then suspend one of the processes

Page-Fault Frequency Scheme



Establish “acceptable” page-fault rate

- If actual rate is too low, remove a frame from that process
- If actual rate is too high, give process a frame
- May need to suspend process if fault rate increases and no free frames available

File Systems

File System Structure

Consists of many levels:

- Logical File System
- File-Organisation Module
- Basic File System
- I/O Control

File attributes

1. Name
2. Type
3. Location
4. Size
5. Time, Date, User ID
 - a. For protection, security, usage monitoring

File Structure

Unstructured:

A file is a stream of bytes. Each byte is individually addressable from beginning of the file

Used by UNIX and MSDOS

File access methods

Sequential Access

Information in the file is processed in order, byte after byte: Operations include:

- Read byte from next position
- Write byte to next position
- Reset to beginning

Direct Access

- Bytes of file can be read in any order by referencing byte number
- Based on disk model of file
- Operations: read byte at n, write byte at n, position to n

File Operations

<i>Commands</i>	<i>Explanation</i>
Create	allocate disk space; create directory entry with file attributes.
Delete	delete the corresponding directory entry; deallocate disk space.
Open	search the directory structure for file entry; move the content to memory (put in the <i>open file table</i>). <u>Information Associated with an open file:</u> - <i>Current Position Pointer</i> - <i>File Open Count</i> - <i>Disk Location</i>
Close	move the content of directory entry in memory to directory structure on disk.
Write	search <i>open file table</i> to find the location of file; write data to the position pointed by <i>Current Position Pointer</i> .
Read	search <i>open file table</i> to find the location of file; read data from the position pointed by <i>Current Position Pointer</i> .

Directory Structure

Directory contains one entry per file, residing on disk

Structures:

1. Each entry contains file name and other attributes
2. Each entry contains file name and pointer to another data structure where attributes are found

Implementation

1. Linear list
 - a. Simple implementation
 - b. Time-consuming to execute, requires linear search
2. Hash table
 - a. Decreases directory search time
 - b. Collisions: two file names hash to same location, use linked list of entries instead


Directory organisation

Logical organisation is needed to have:

1. Efficiency of file locating
2. Naming
 - a. Different users, same name for different files
 - b. Same file, different names
3. Grouping
 - a. Logical grouping of files by their properties

Two Level Directory

One directory for each user:



Efficiency: Pass

Naming: Pass

Grouping: Fail

Tree-Structured Directories

- Path name:
 - Absolute: Begins at root and follows path down to specific file
 - Relative: Defines path from current directory
- Characteristics
 - Efficiency: Pass
 - Naming: Pass
 - Grouping: Pass

Acyclic-graph Directory

Graph with no cycles is natural generalisation of the tree scheme

Naming: File can have two different names (aliasing)

Support for File Sharing: Necessary for collaboration

File Sharing

1. Symbolic Link
 - a. Directory entry is link, containing absolute/relative path name
 - b. Resolve link by using path name to locate real file
 - c. Slower access than hard linking
2. Hard link
 - a. Duplicate all information about file (File Control Block) in multiple directories

Issues with File Sharing:

- In traversing the file system, shared files visited more than once
 - Need to ignore the link
- Deleting a shared share leaves dangling pointers to non-existent file
 - Solutions:
 - Search for dangling links and remove them
 - Leave dangling links and delete them only when they are used again
 - Preserve the file until all references are deleted

Directory operations

<i>Commands</i>	<i>Explanation</i>
Create	create a directory. In UNIX, two entries "." and ".." are automatically added when a directory is created. "." refers to the <i>current directory</i> ; and ".." refers to its <i>parent</i> .
Delete	delete a directory. Only empty directory can be deleted (directory containing only "." and ".." is considered empty).
List	list all files (directories) and their contents of the directory entry in a directory
Search	search directory structure to find the entry for a particular file.
Traverse	access every directory and every file within a directory structure.

File protection in UNIX

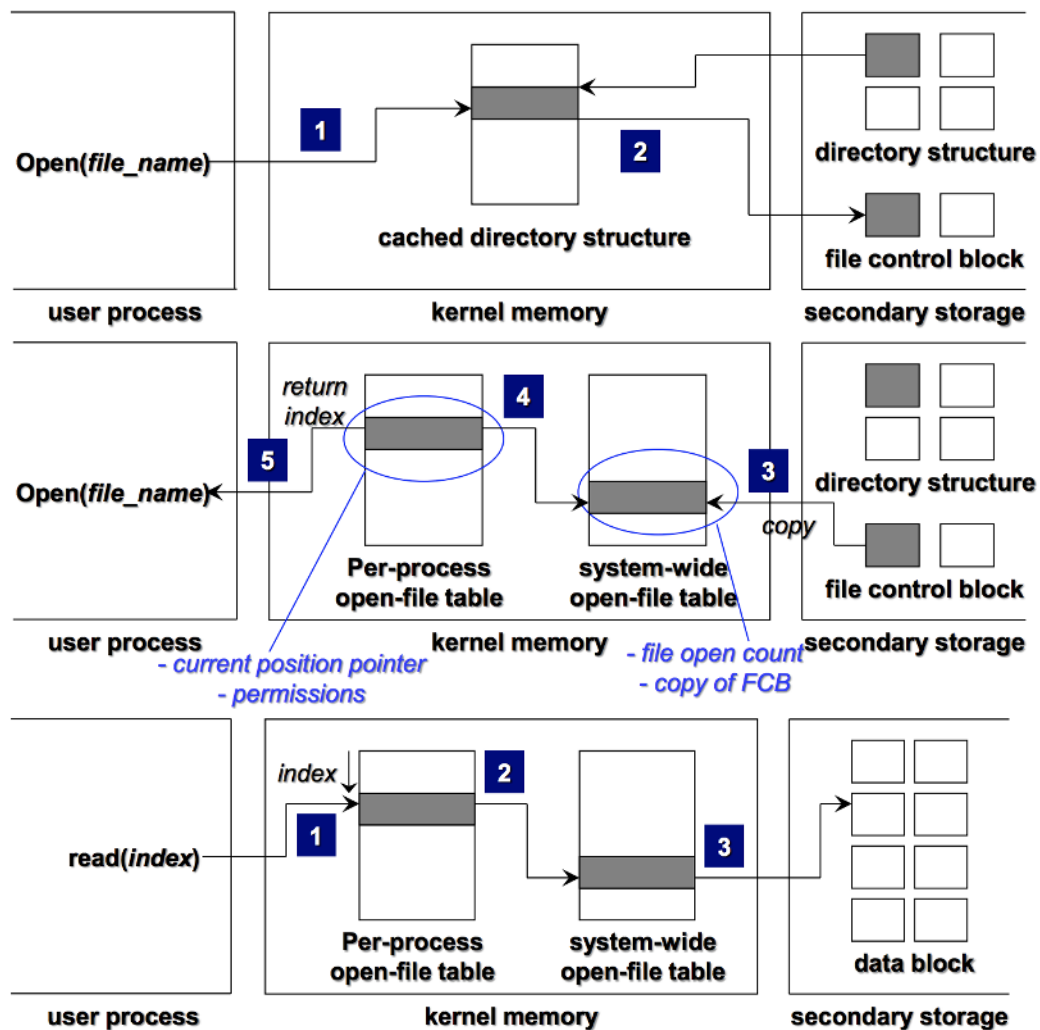
Model of Access: Read, Write, Execute

Three classes of users: Owner, Group and Public access

Permissions on a directory:

- To access a directory, the execute permission is needed. Without it, one cannot execute any command on that directory or have access to file contained in the file hierarchy rooted at that directory
- No read permission: Cannot list directory
- No write permission: Can't create or delete files in directory

In-memory File System Data Structure




Allocation Methods

1. Contiguous allocation
2. Linked allocation
3. Indexed Allocation

Contiguous allocation

Each file occupies set of contiguous blocks on disk

- Simple: starting location (block #) and length (number of blocks) is required
 - Supports random access
- Problems: Waste of space (fragmentation etc)
 - Find space constricted by size of hole, may need to move to bigger hole
 - If needed file space is overestimated: Internal fragmentation
- Logical to physical address mapping:
 - Block size 512 bytes
 - Logical address $Q \times 512 + R$
 - Block to be accessed = $Q + \text{starting address}$
 - Displacement into block = R



directory		
file	start	length
count	0	2
tr	14	3
mail	19	6
list	28	4
f	6	2

Logical Address = 2333

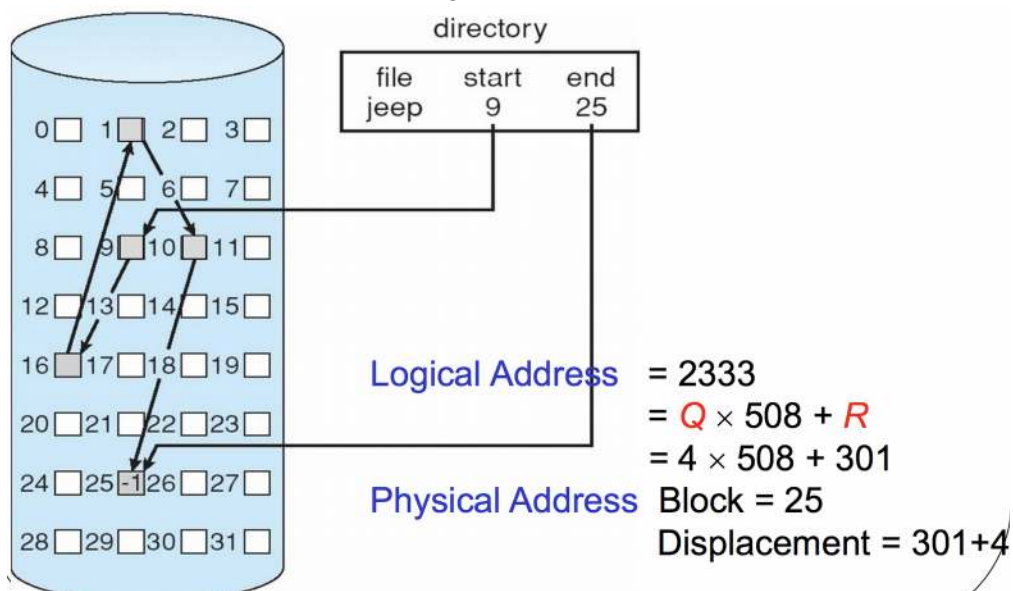
$$= Q \times 512 + R$$
$$= 4 \times 512 + 285$$

Physical Address = 285th byte of block
23

Linked Allocation

Each file is a linked list of disk blocks, blocks are scattered anywhere on disk

- Simple: need only starting address
 - No waste of space
 - No constraints on file size: blocks allocated as needed
- Problem
 - Random Access not supported
- Logical to Physical Address Mapping:
 - 512 bytes, first 4 bytes reserved for pointer to next block in the list
 - Logical address: $Q \times 508 + R$
 - Block to be accessed is (Q+1)th block in linked chain of blocks representing file
 - Displacement into block = $R+4$
- Allocate as needed, link together



Linked Allocation: File Allocation Table


Disk space allocation used by MS-DOS and OS/2

- File Allocation Table has one entry for each disk block, indexed by block number
- The FAT entry contains either a special end-of-file value, or block number of next block in file, or a "free"
- Directory entry contains block number of first block
- Random access can be optimised

Indexed Allocation

Each file has index block containing pointer to allocated blocks. Directory entry contains block number of index lock

- Supports random access
- Dynamic storage allocation without external fragmentation
- Problem:
 - Overhead of keeping index blocks and address mapping
- Logical to physical address mapping:
 - Maximum file size is 128K bytes and block size is 512 bytes
 - 2 blocks needed for index table (4 bytes per pointer)
 - Logical Address: $Q \times 512 + R$
 - Displacement into index table: Q
 - Displacement into block: R



Index Allocation: UNIX inode

- For each file/directory, there is an inode (index block)
- Inode contains:
 - File attributes
 - 12 pointers to direct data blocks
 - 3 pointers to indirect index blocks
 - Single indirect
 - Double indirect
 - Triple indirect

With 4-byte file pointer and 4K bytes block.

Max file size:

Direct access $12 \times 4K = 48K$

Single indirect access: $4K \times 4K/4 = 2^{22}$

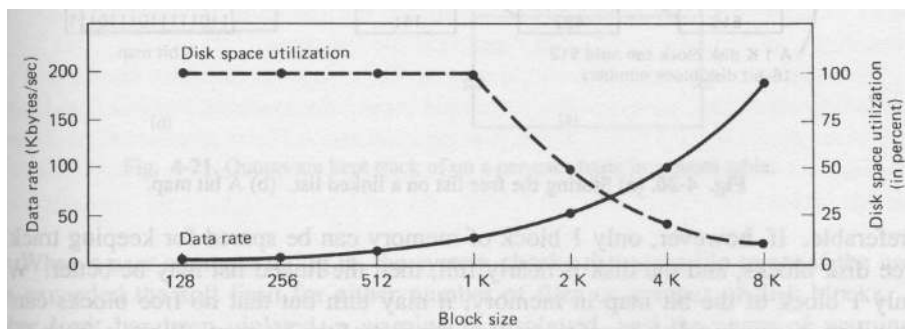
Double indirect access $4K \times 4K/4 \times 4K/4 = 2^{32}$

Total file size: > 4GB

Disk-Space Management

Block size affects both data rate and disk space utilisation

- Big block size: File fits into few blocks: Fast to find and transfer, wastes space if file does not occupy entire last block
- Small block size: File consist of many blocks: Slow data rate
- Trade-off between time and space utilisation



Keeping track of free blocks

Bitmap or Bit Vector

Linked List

Bitmap

Block size: 2^9

Disk Size: 2^{34} (16 GB)

Bit map size = $(2^{34}/2^9)/8 = 2^{22}$ bytes


- Bit map usually kept in fixed place on disk, brought into memory for efficiency, write back to disk occasionally for consistency and security
- Easy to locate free blocks but inefficient unless entire map is kept in memory

Linked list

Link all free disk blocks together, keeping pointer to first free block in special location on disk and caching it in memory

No extra space required but not efficient

I/O System and Disk



I/O System Design Objectives

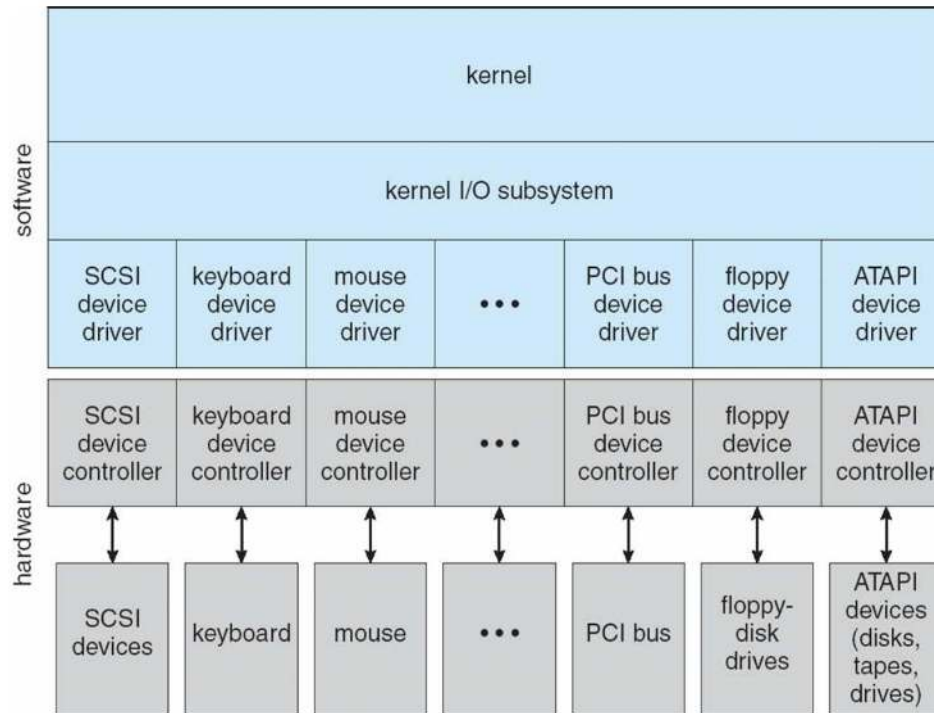
Efficiency

- I/O operations often form bottleneck
- I/O operations are extremely slow compared with main memory and processor
- Area with most attention is disk I/O

Generality

- Desirable to handle all devices in uniform manner
- Applies to the way processes view I/O devices and OS manages I/O devices and operations

Kernel I/O Structure





Kernel I/O Subsystem

- **I/O Scheduling**
 - Schedules I/O request through rearranging the order of service
 - Can improve system efficiency
- **Buffering**
 - Store data in memory while transferring devices
 - Cope with device speed mismatch
 - Cope with device transfer size mismatch like in networking
- **Caching**
 - Fast memory holding copy of data
 - Improve I/O efficiency for files written and reread rapidly
- **Spooling**
 - Hold output for device
 - Used in device serving only one request at a time like a printer

Performance Consideration

1. Synchronous/ Blocking I/O →
2. Asynchronous/ Non-blocking



asynchronous. CPU issues request and continues running other processes.

Data copying between I/O controllers and physical memory also affect system performance

Disk Scheduling

OS is responsible for using hardware efficiently: Disk drives must have fast access time and high disk bandwidth

Access time of a disk:

1. Seek time: Disk to move heads to track containing desired sector (5-25 ms)
2. Rotational latency: Disk rotate desired sector to disk head (8ms)

Aim: Minimise seek time (seek distance)

Disk bandwidth is total number of bytes transferred, divided by total time between 1st request and completion of last transfer

Disk Scheduling Algorithms

FCFS

FCFS may have wild swings back and forth

Shortest-Seek-Time-First (SSTF)

Selects request with minimum seek time from current head position

SSTF is a form of SJF: Starvation under heavy load as distant requests may never be serviced

Scan/Elevator algorithm

Disk arm starts at one end of disk, moves towards other end, servicing requests until it gets to other end where head movement reversed and servicing continues

Unlikely Starvation but requests may be delayed

Circular-SCAN (C-SCAN)

Head moves from low end to high end of disk, servicing requests along

When reaching the high end, head immediately back to beginning of disk without servicing

C-LOOK

Moves from low end to high end

Reaches last request, head returns immediately to lowest cylinder # request

Disk servicing influenced by file-allocation method:

Program reading contiguously allocated file generate requests close together

Linked or indexed file generates requests widely apart resulting in greater head movement

Selecting a disk-scheduling algorithm

The algorithm should be written as a separate module, allowing it to be replaced with different algorithm if necessary.

SSTF or C-LOOK is a reasonable choice.

Disk Management

Low-level formatting (physical formatting) Dividing a disk into sectors that the disk controller can read and write

To use a disk to hold files, the OS needs to record its own data structure on the disk

1. Partition the disk into one or more groups of cylinders
2. Logical formatting or "making a file system"
3. Bootstrap program initializes system
 - a. Tiny bootstrap stored in ROM, loads full bootstrap program (in the disk boot

- blocks) into memory and starts execution
- b. Finds OS kernel on disk and loads into memory
- c. Jumps to initial address to start OS

Disk Reliability

Disk striping uses group of disks as one storage unit

Each block is broken into sub-blocks, each sub-block stored on each disk

Block transfer is faster due to sub-blocks transferred in parallel

Disk mirroring keeps duplicate of each disk

Logical disk consisting of 2 physical disks

If one fails, the data can be read from the other

RAID (Redundant-Array-of-Independent-Disks)

RAID 0: Non-redundant striping

RAID 1: Mirror duplicates

Mirror of Strips: RAID 0 + 1

Strips of Mirrors: RAID 1 + 0