

Quiz 01

What is the difference between a long-term and medium-term scheduler?

- ☐ Long-term scheduler schedules processes in main memory on the CPU; Medium-term scheduler decides which active processes to keep in main memory and which ones to swap to disk when load is heavy.
- ☐ They are identical.
- ☐ Long-term scheduler decides which new processes should be brought to main memory from disk; Medium-term scheduler schedules processes in main memory on the CPU.
- ☒ Long-term scheduler decides which new processes should be brought to main memory from disk; Medium-term scheduler decides which active processes to keep in main memory and which ones to swap to disk when load is heavy.

What is the difference between a process and a program?

- ☐ They are identical
- ☐ Every user program can have at most one active process in the system.
- ☐ They are two different and unrelated concepts.
- ☒ Process is an instance of a program that is currently active in the system

What are the typical resource constraints of embedded systems?

- ☐ Small memory and low power (battery operated).
- ☐ Low network bandwidth and small displays.
- ☐ Slow processors
- ☒ All the other options are correct.

What is an Operating System?

- ☐ Software program that controls other user programs.
- ☐ Allocates and manages hardware resources.
- ☐ Always ready to accept new commands from users and hardware.
- ☒ All the other options are correct.

Multiprogrammed systems allow parallel execution of jobs, hence they are also called time-shared systems.

- ☒ False: Parallel execution of jobs depends on the hardware (multiprocessing); multiprogramming allows efficient switching between jobs for CPU access (time-division access), and hence it is also called time-shared system.
- ☐ True: Multiprogramming allows simultaneous access to the CPU for multiple user jobs, and hence it is also called time-shared system.

Kernel or monitor or supervisor mode is the same as "superuser" in Linux.

- ☒ False: Kernel mode is a hardware mode of operation, whereas "superuser" is the "root" user account.
- ☐ True: Both, the kernel mode as well as "superuser" allows one to execute privileged instructions in hardware.
- ☐ False: "superuser" is more privileged than kernel mode.
- ☐ False: "superuser" is the same as supervisor mode, whereas kernel or monitor mode is more privileged.

What parts of the process memory space are shared between threads in the same process.

- ☐ Everything is shared between threads in a single process.
- ☐ Only Text and OS resources (files, etc.) are shared; all other parts are private to the threads.
- ☒ Text, Data, Heap and OS resources (files, etc.) are shared between the threads; Stack and Registers are not.
- ☐ Text, Data and Heap are shared between the threads; OS resources (files, etc.), Stack and Registers are not.

Quiz 01

What are the four memory regions of a process and how are they used.

- ☐ "Text" for documentation about the process; "Data" for all variables; "Stack" for managing function calls; "Heap" for garbage collection.
- ☐ "Code" for instructions; "Global" for static and global variables; "Function" for managing function calls and local function variables; "Dynamic" for dynamically created variables.
- ☒ "Text" for instructions; "Data" for static and global variables; "Stack" for managing function calls and local function variables; "Heap" for dynamically created variables.
- ☐ All the other options are correct.

In real-time systems, average response time of user programs should be minimized.

- ☐ True: This will ensure that the required deadlines are met.
- ☐ True: This will ensure that the system is responsive.
- ☒ False: Real-time systems require guaranteed satisfaction of deadlines, irrespective of average response time.
- ☐ False: Real-time systems require average response times to be maximized.

What is the difference between shared memory and message passing based Inter-Process Communication.

- ☐ They are two different names for the same communication mechanism which uses kernel memory space.
- ☒ In message passing, the communication is handled by the OS; In shared memory, the communication is handled directly by the processes through a common memory space.
- ☐ In shared memory, there is memory in the kernel space that is shared between the processes; In message passing, no memory is required in the kernel memory space.
- ☐ They are two different names for the same communication mechanism which uses user memory space.

What is an advantage and a disadvantage of many-to-many mappings of logical (user) to physical (kernel) threads?

- ☐ High concurrency and mapping is very simple many kernel threads would be required.
- ☒ High concurrency and not many kernel threads required; mapping is very complex.
- ☐ Mapping is very simple and not many kernel threads required; concurrency is low.
- ☐ High concurrency; mapping is very complex and many kernel threads would be required.

What is the disadvantage of many-to-one mapping between logical (user) and physical (kernel) threads?

- ☐ The mapping process is very complex.
- ☒ A blocked user thread will block all the other user threads that are mapped to the same kernel thread.
- ☐ There is no disadvantage.
- ☐ A blocked user thread will also block all the other kernel threads in the system.

DMA does not use interrupts.

- ☒ False: DMA uses interrupts once the block transfer is complete (between memory and I/O device) to notify the CPU.
- ☐ True: DMA bypasses the CPU and hence does not require interrupts.

Describe one advantage and one disadvantage of batch systems.

- ☐ Memory management is simple (1 user job); CPU utilization is not efficient (scheduling of jobs is complex).
- ☒ Memory management is simple (1 user job); CPU utilization is not efficient (idling during I/O).
- ☐ CPU utilization is very efficient (1 user job to schedule); Memory management is complex (deciding which user job to keep in memory).

What comprises a process context and where is it stored?

- ☒ All registers required to execute/resume the process (program counter, stack pointer, memory protection registers, etc.) and it is stored in the PCB.
- ☐ The process memory as defined by the base and limit register values and it is stored in the user memory space.
- ☐ All registers required to resume the process (program counter, stack pointer, memory protection registers, etc.) and it is stored in the user memory space.
- ☐ Only base and limit register values for a process and it is stored in the PCB.

Quiz 01

Which of the following precisely describes the sequence of steps to process an interrupt.

- ☐ Disable interrupts, locate interrupt service routine (ISR) in interrupt vector table, execute ISR, enable interrupts.
- ☒ Switch to kernel mode, save program context, locate interrupt service routine (ISR) in interrupt vector table, execute ISR, restore program context, switch to user mode.
- ☐ Switch to kernel mode, locate interrupt service routine (ISR) in interrupt vector table, execute ISR, switch to user mode.
- ☐ Disable interrupts, save program context, locate interrupt service routine (ISR) in interrupt vector table, execute ISR, restore program context, enable interrupts.

What is the difference between multi-threading, multi-programming and multi-processing?

- ☐ They are all different terms for the same concept of time-multiplexed execution.
- ☐ multi-threading and multi-programming are identical; multi-processing means multiple CPU cores for execution.
- ☐ multi-threading means multiple threads of execution within a process; multi-programming and multi-processing are identical.
- ☒ multiple threads of execution within a process; multiple processes ready for execution in main memory; multiple CPU cores for execution.

What are the key OS features needed for multiprogramming.

- ☐ Multiprocessing, DMA and I/O protection.
- ☐ CPU scheduling, multiprocessing and DMA.
- ☒ I/O device management, CPU scheduling and memory management.
- ☐ Memory management, I/O protection and CPU scheduling.

There is one base and one limit register in the hardware for each user program.

- ☐ True: This is necessary to ensure memory protection for each user program.
- ☒ False: There is only one base and one limit register overall in the hardware, which are loaded with appropriate values by the OS before executing the user program.
- ☐ False: There is only one base and one limit register overall in the hardware, which are loaded with appropriate values by the CPU before executing the user program.
- ☐ False: These registers are in the OS, which also checks whether each memory access is within range.

Trap is an interrupt generated only due to software exceptions.

- ☐ True: CPU transfers control to the OS to process the exception.
- ☒ False: CPU generates trap either due to software exception or system call, and transfers control to the OS.
- ☐ False: OS generates trap due to software exceptions and transfers control to the CPU.
- ☐ False: OS generates trap and processes it.

Multiprocessing without multiprogramming is efficient?

- ☐ Yes, they are different techniques that can applied independently and all combinations are efficient.
- ☐ Yes, even without multiprogramming multiprocessing will ensure high CPU utilization.
- ☒ No, without multiprogramming there will be only one user job to execute and hence CPU utilization will be low.

System calls are identical to library function calls in a high-level programming language such as C.

- ☐ True: For example, fopen() in C programming is a system call to open a file.
- ☒ False: User programs can access services in the OS by invoking software functions implemented in the OS through system calls.

When a user program tries to execute an I/O instruction directly,

- ☐ the CPU executes the instruction
- ☐ the CPU checks if the request is within range of the base and limit registers
- ☐ the CPU checks for its validity and legality before executing it.
- ☒ the CPU generates a trap and transfers control to the OS.