

1 Objective

In this group assignment, you will build an **opinion search engine**. Given a specific topic of your choice (e.g., cryptocurrencies), your system should enable users to **find relevant opinions about any instance of such topic** (e.g., bitcoin) and **perform sentiment analysis on the results** (e.g., opinions about bitcoin are 70% **positive** and 30% **negative**). Once you have chosen a topic, make sure that a) you can **find enough data about it** (e.g., some topics may be too niche to the point that you would only find a few hundreds data about it) and b) the **opinions you get are balanced** (e.g., if the topic you chose only has negative opinions associated with it, it's probably not a good topic). For ideas about interesting topics, you can check our project page at <https://sentic.net/projects>. You are pretty much free to use anything you want in terms of available tools/libraries. However, your system cannot be just a mashup of existing services. Your **final score** will depend not only on how you developed your system but also on its **novelty and your creativity**: in other words, to get a high score you do not only need to **implement a system that works**, but also a **system that is useful and user-friendly**.

2 Deadline

The assignment constitutes 35% of your total grade for the course. Assignments are to be submitted via Blackboard (email submissions will not be considered) by **end of week 12th (Sunday at 11:59pm SGT)**. **5% points will be deducted for each rounded-off day after the deadline**. **Only the first submission counts** (Blackboard allows for multiple submissions only in case of system errors or disconnections).

3 Grouping

The assignment will be done in groups of **6 or 5 people**. Members will be randomly assigned by the system for fairness. The main tasks of the assignment are **crawling (20 points)**, **indexing (40 points)** and **classification (40 points)**. If you like, you can split your group into up to **three subgroups** taking care of each of these tasks and **specify who did what in your final report** so that each member will be graded accordingly. If this information is not specified, a unique grade will be given to the whole project and this will be shared among all members of the group (recommended option). Some **overlap between projects from different groups is allowed** but beware that, if we find out that a project has more than 30% overlap with another project from this year or past years, your group will be disqualified (and get zero points as final grade for the assignment). Hence, it is **OK to share the general idea of your assignment with other groups but not the implementation details**.

4 Tasks and Questions

4.1 Crawling (20 points)

Crawl **text data from any sources** which you are interested in and permitted to access, e.g., **Twitter API or Reddit API**. The crawled corpus should have at least **10,000 records and at least 100,000 words**. It is **OK to use available datasets for training** (e.g., **popular sentiment benchmarks**), but you still have to at least crawl and label data for testing. For your own evaluation dataset, **use the same format as the above-mentioned sentiment benchmarks** (using a different format will result in demerit points) and name it **"eval.xls"**. Also, make sure your **dataset does not contain duplicates** and try your best to **make it balanced** (e.g., **equal number of positive and negative entries**). Before crawling any data, carefully consider the questions in this material, e.g., check whether the **data have enough details to answer the questions**. You can **use any third party libraries** for the crawling task, e.g.:

- Jsoup: <https://jsoup.org>
- Twitter4j: <https://twitter4j.org>
- Facebook marketing: <https://developers.facebook.com/docs/marketing-apis>
- ✗ Instagram: <https://instagram.com/developer>
- Amazon: <https://github.com/ivanpgs/amazon-crawler>
- Tinder: <https://gist.github.com/rtt/10403467>
- Tik Tok: <https://developers.tiktok.com>

Question 1: Explain and provide the following:

1. How you crawled the corpus (e.g., source, keywords, API, library) and stored it
2. What kind of information users might like to retrieve from your crawled corpus (i.e., applications), with sample queries
3. The numbers of records, words, and types (i.e., unique words) in the corpus

4.2 Indexing (40 points)

Indexing: You can do this from scratch or use some combination of available tools, e.g., Solr+Lucene+Jetty. Solr runs as a standalone full-text search server within a servlet container such as Jetty. Solr uses the Lucene search library at its core for full-text indexing and search, and has REST-like HTTP/XML and JSON APIs that make it easy to use from virtually any programming language. Useful documentations include:

- Solr project: <https://lucene.apache.org/solr>
- Solr wiki: <https://wiki.apache.org/solr/FrontPage>
- Lucene tutorial: <https://lucene.apache.org/core/quickstart.html>
- Solr with Jetty: <https://wiki.apache.org/solr/SolrJetty>
- Jetty tutorial: <https://wiki.eclipse.org/Jetty/Tutorial>

You can also choose other inverted-index text search engine open projects, e.g., Sphinx, Nutch, and Lemur. However, you should NOT simply adopt SQL-based solutions for text search (for example, you CANNOT solve text search simply using Microsoft Sqlserver or MySQL).

Querying: You need to provide a simple but friendly user interface (UI) for querying. It could be either a web-based or mobile app based UI. You could use JSP in Java or Django in Python to develop your UI website. Since Solr provides REST-like APIs to access indexes, one extra JSon or RESTful library would be enough. Otherwise, you may use any third party library. The UI must be kept simple. A sophisticated UI is not necessary nor encouraged (as it is not the focus of this course). Detailed information besides text is allowed to be shown for the query results, e.g., product images on Amazon, ratings on Amazon, and pictures on Instagram. The details should be designed to solve specific problems.

Question 2: Perform the following tasks:

- Design a simple UI (you can design one from scratch or you can tap on an existing one, e.g., Solr UI) to allow users to access your system in a simple way
- Write five queries, get their results, and measure the speed of the querying

Question 3: Explore some innovations for enhancing the indexing and ranking. Explain why they are important to solve specific problems, illustrated with examples. You can list anything that has helped improving your system from the first version to the last one, plus queries that did not work earlier but now work because of the improvements you made. Possible innovations include (but are not limited to) the following:

- ☒ Timeline search (e.g., allow user to search within specific time windows)
 - Geo-spatial search (e.g., use map information to refine query results and improve visualization)
- ☒ Enhanced search (e.g., add histograms, pie charts, word clouds, etc.)
 - Interactive search (e.g., refine search results based on users' relevance feedback)
- ☒ Multimodal search (e.g., implement image or video retrieval)
- ☒ Multilingual search (e.g., enable your system to retrieve data in multiple languages)
- ☒ Multifaceted search (e.g., visualize information according to different categories)

inverted index for different sections of
scraped excel data

4.3 Classification (40 points)

Although often defined as a **binary categorization problem**, sentiment analysis is actually a complex task, or **suitcase research problem**, as it **requires tackling many other subtasks**. Choose **at least two subtasks to perform information extraction on your crawled data**. Unless you are sure that your **data does not contain any neutral content**, you should always cover at least **subjectivity detection and polarity detection**. Namely, you should first categorize your data as **neutral** versus **opinionated** and then **classify** the resulting opinionated data as **positive** versus **negative**. Different classification approaches can be applied, including:

- knowledge based, e.g., SenticNet
- rule based, e.g., linguistic patterns
- machine learning based, e.g., deep neural networks
- hybrid (a combination of any of the above)

You can tap into any resource or toolkit you like, as long as you motivate your choices and you are able to critically analyze obtained results. Some possible choices include:

- Weka: <https://cs.waikato.ac.nz/ml/weka>
- Hadoop: <https://hadoop.apache.org>
- ✕ Pylearn2: <https://deeplearning.net/software/pylearn2>
- SciKit: <https://scikit-learn.org>
- ✕ NLTK: <https://nltk.org>
- Theano: <https://github.com/Theano>
- Keras: <https://github.com/fchollet/keras>
- Tensorflow: <https://github.com/tensorflow/tensorflow>
- PyTorch: <https://pytorch.org>
- Huggingface: <https://huggingface.co/>
- AllenNLP <https://github.com/allenai/allennlp>

Question 4: Perform the following tasks:

- **Motivate** the **choice of your classification approach** in relation with the state of the art
- **Discuss whether you had to preprocess data** (e.g., microtext normalization) and **why**
- Build an **evaluation dataset by manually labeling 10% of the collected data** (at least 1,000 records) **with an inter-annotator agreement of at least 80%** (it is recommended to have **3 annotators**, but 2 is also OK)
- **Provide evaluation metrics** such as **precision, recall, and F-measure** on such dataset
- **Perform a random accuracy test** on the **rest of the data** and **discuss results**
- **Discuss performance metrics**, e.g., records classified per second, and scalability of the system

Question 5: Explore some **innovations for enhancing classification**. If you introduce more than one, perform an **ablation study to show the contribution of each innovation**. For example, if you perform word sense disambiguation (WSD) and named entity recognition (NER) to enhance sentiment analysis, **show the increase in accuracy when adding only WSD**, the increase in accuracy when adding only NER, and the **increase in accuracy when adding both** WSD and NER to your system. **Explain why they are important to solve specific problems, illustrated with examples**. **Possible innovations include** (but are not limited to) the following:

- Enhanced classification (add another sentiment analysis subtask, e.g., sarcasm detection)
- Fine-grained classification (e.g., perform aspect-based sentiment analysis)
- ? • Hybrid classification (e.g., apply both symbolic and subsymbolic AI)
- Cognitive classification (e.g., use brain-inspired algorithms)
- Multitask classification (e.g., perform two sentiment analysis tasks jointly)
- Ensemble classification (e.g., use stacked ensemble)

5 Submission

Select one person of the group in charge of submitting the final assignment report. Submission has to be done via Blackboard (do not email your report). The results will not be presented in person so make sure your report is clear and self-contained. However, you do have the chance to present your work through a short YouTube video. In fact, if you search CZ4034 on YouTube, you can view some examples from past years. You can take inspiration from them but doing a project that is more than 30% similar to past (and current) ones will result in the annulment of your assignment. The submission shall consist of one single PDF file named after your group number, e.g., if you are group 10, your file should be titled simply “10.pdf”. Failing to name the file correctly or sending it in the wrong format, e.g., zip or MS Word, will result in demerit points. Do add some pictures to your report to make it clearer and easier to read. There is no page limit and no special formatting is required (but please use the new NTU logo instead of the old one). The file shall contain the following five key items:

1. The names of the group members + matriculation number in the first page
2. Your answers for all the above questions
3. A YouTube link to a video presentation of up to 5 minutes; in the video, introduce your group members and their roles, explain the applications and the impact of your work and highlight, if any, the creative parts of your work (note that you do not have to give all the answers in the video presentation)
4. A Dropbox (or similar, e.g., Google Drive or OneDrive) link to a compressed (e.g., zip) file with crawled text data, queries and their results, evaluation dataset, automatic classification results, and any other data for Questions 3 and 5
5. A Dropbox (or similar, e.g., Google Drive or OneDrive) link to a compressed (e.g., zip) file with all your source codes and libraries, with a README file that explains how to compile and run the source codes

Good luck! :)